

Convolutional Neural Networks and Data Augmentation

June course: Deep Learning in Computer Vision

Morten Hannemose, mohan@dtu.dk

June 6th, 2019


Outline – What you're going to see

- Convolutions recap
- CNN
 - Convolutions
 - Max pooling
 - Stride/padding
 - Backprop
- Combatting overfitting:
 - Dropout
 - Data augmentation


What is a CNN?

- Convolutional Neural Network
 - Uses convolutions
- Local connectivity
- Weight sharing

Before:
(Fully connected)



After:
(Convolutional)



Convolutions

- Convolutions and cross correlation are related operations, but convolution involves rotating the filter 180 degrees.
 - We will refer to cross correlation as “convolution” in this course, as is common in deep learning

9	1	2	2	7
10	3	10	5	1
2	6	10	10	9
10	10	5	8	10
7	10	9	10	7

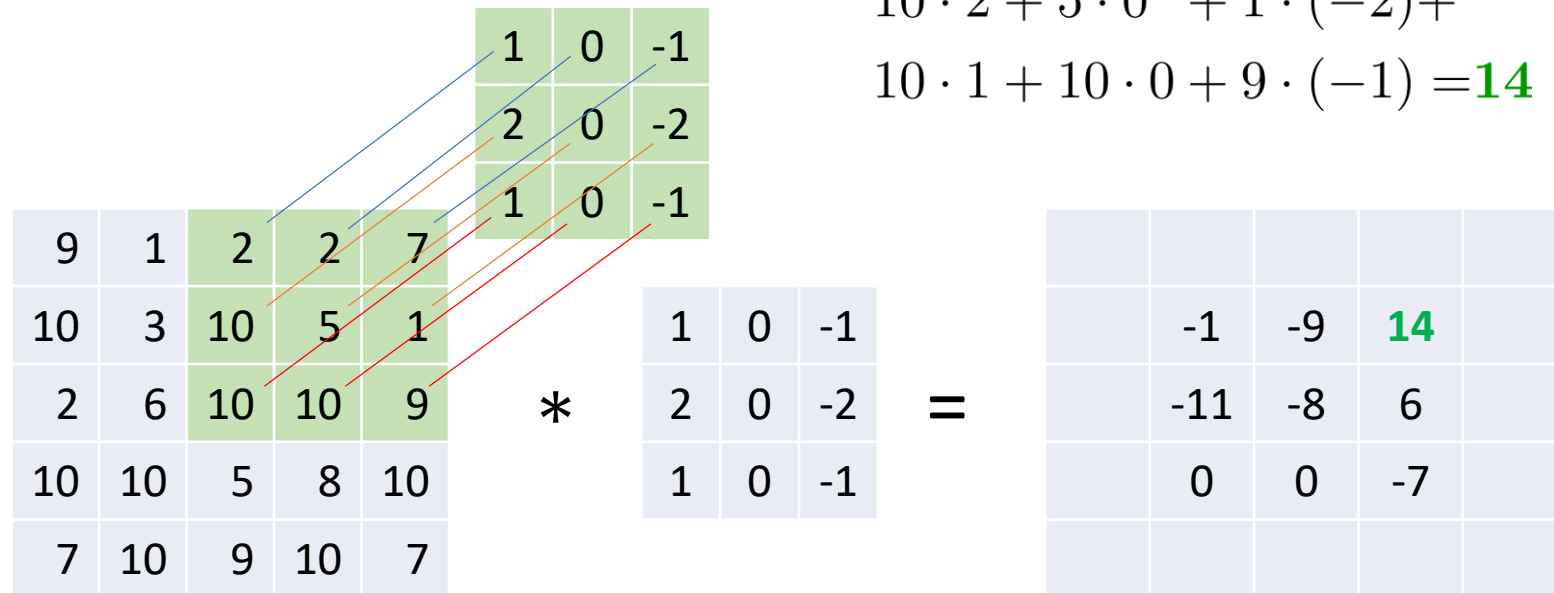
 $*$

1	0	-1
2	0	-2
1	0	-1

 $=$

	-1	-9	?	
	-11	-8	6	
	0	0	-7	

Convolutions



Convolutions

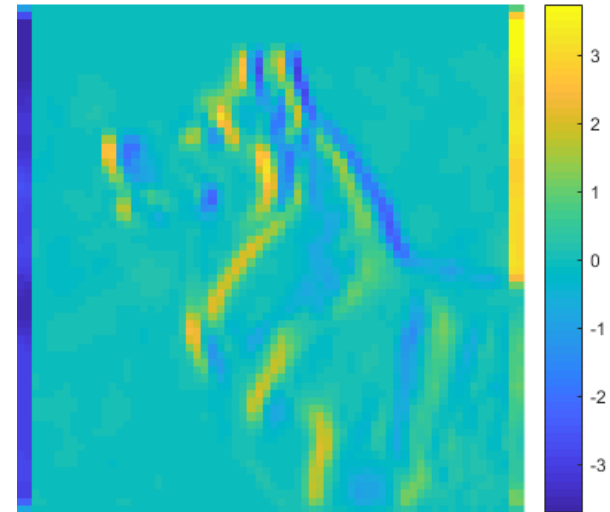
- This filter detects vertical edges



*


1	0	-1
2	0	-2
1	0	-1

=



Convolutions

- What do we do when images have multiple channels?
e.g. color images



*

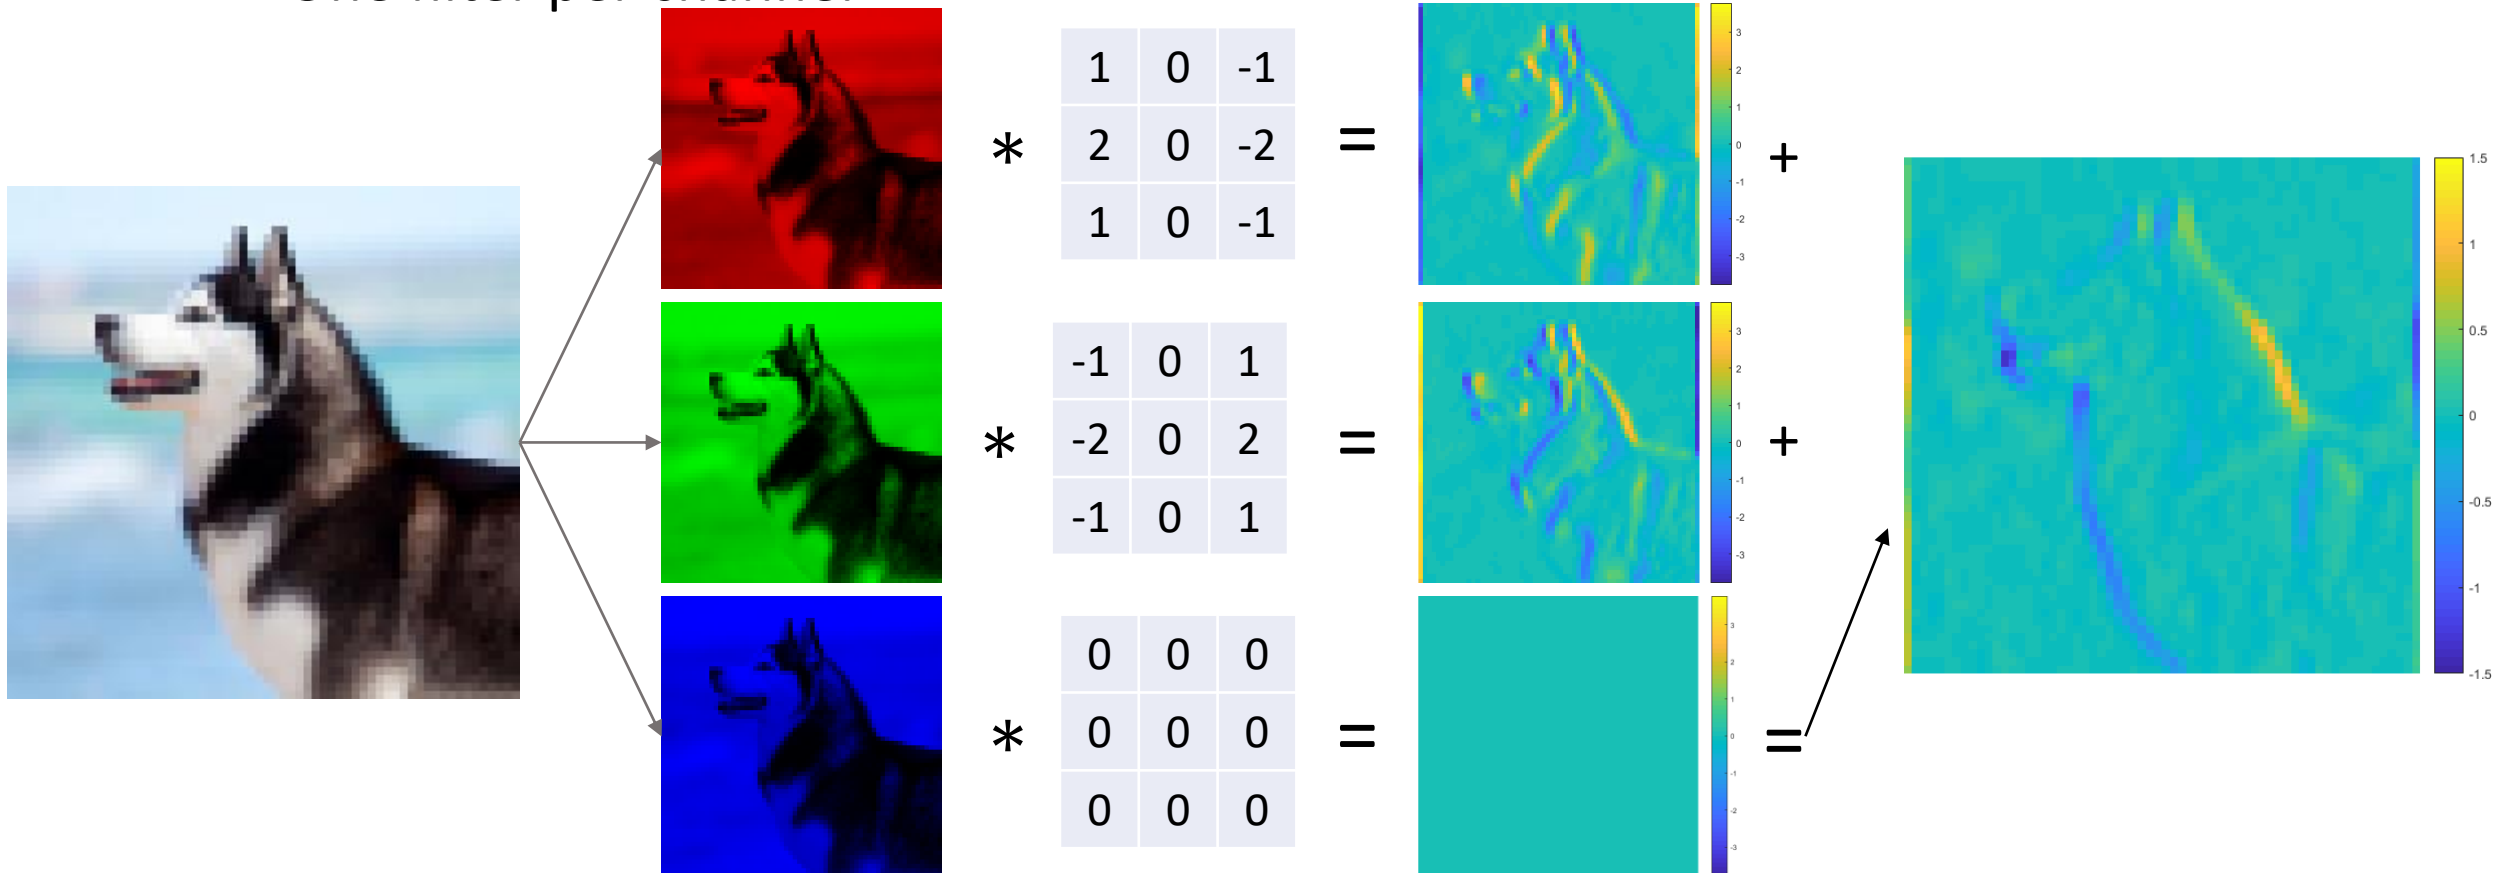
1	0	-1
2	0	-2
1	0	-1

=




Convolutions

- What do we do when images have multiple channels?
 - One filter per channel




What is a CNN?

- Exploit 2d layout of images
- Images are volumes
- Color image → three channels
 - Represented in computer as
 $3 \times 64 \times 64$




What is a CNN?

- Let's look at a single neuron
- Input is 3 channels, 64×64
- 7×7 convolution
- Each neuron looks at a $3 \times 7 \times 7$ volume in the input layer
 - Requires as many weights + one bias
- Spatially: locally connected
- Depthwise: Fully connected



What is a CNN?

- Multiple neurons:
- 6 channels (features) output
- Each layer in the output has its own weights
- $6 \times 3 \times 7 \times 7$ weights for this layer



Mathematical definitions

- Forward pass:

$$z_{i,j}^{\ell,f} = \sum_{c=1}^C \sum_{m=-M}^M \sum_{n=-N}^N w_{m,n}^{\ell,f,c} \cdot a_{i+m,j+n}^{\ell-1,c}$$



$$a_{i,j}^{\ell,f} = \sigma \left(z_{i,j}^{\ell,f} \right)$$

$$\delta_{i,j}^{\ell,f} = \frac{\partial \mathcal{L}}{\partial z_{i,j}^{\ell,f}}$$

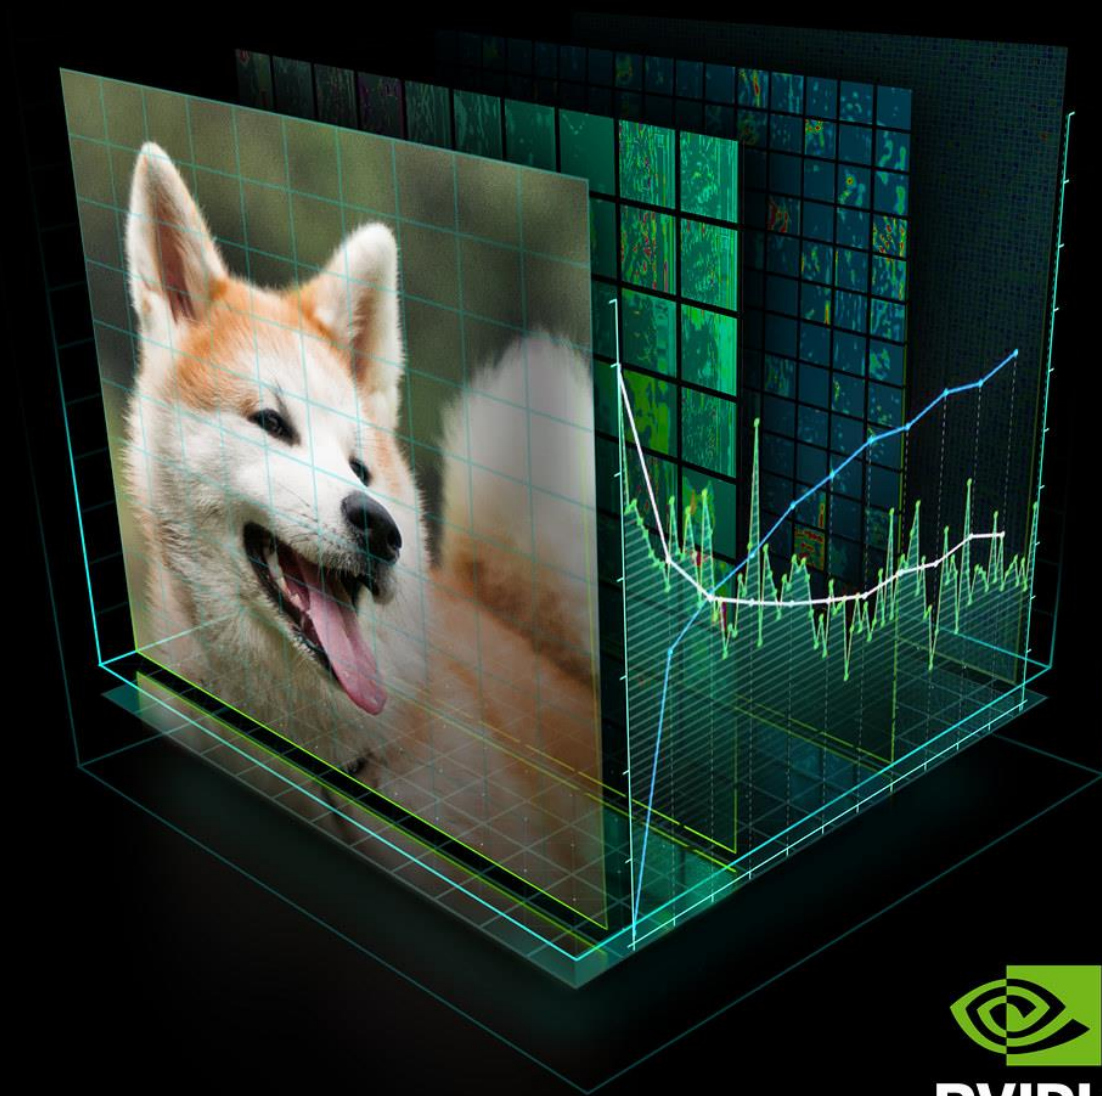
- Backward pass:

$$\delta_{i,j}^{\ell,c} = \frac{\partial \mathcal{L}}{\partial z_{i,j}^{\ell,f}} = \sigma' \left(z_{i,j}^{\ell,c} \right) \sum_{f=1}^F \sum_{m=-M}^M \sum_{n=-N}^N w_{m,n}^{\ell+1,f,c} \cdot \delta_{i-m,j-n}^{\ell+1,c}$$

$$\frac{\partial \mathcal{L}}{\partial w_{m,n}^{\ell,f,c}} = \sum_i \sum_j \delta_{i,j}^{\ell,f} \cdot a_{i+m,j+n}^{\ell-1,c}$$




Wow such image!



nvidia.

What is a CNN?

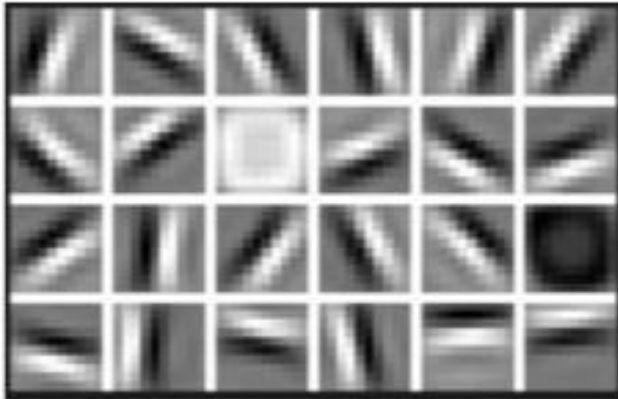
- Local connectivity
- Weight sharing
- Usually followed by a fully connected network to output a classification



Minibatches revisited

- How do we store a minibatch of images in the computer?
 - 4d tensor with size
 - NCHW (minibatch dimension, channels, height, width)
 - Tensorflow has default (which is slower)
 - NHWC
- Minibatches should be made of data sampled without replacement from your full dataset
 - Once all data has been show to the network once it is called an **EPOCH**
 - After an *epoch* you start sampling all your data over again.

Intuition



First Layer Representation



Second Layer Representation




Third Layer Representation

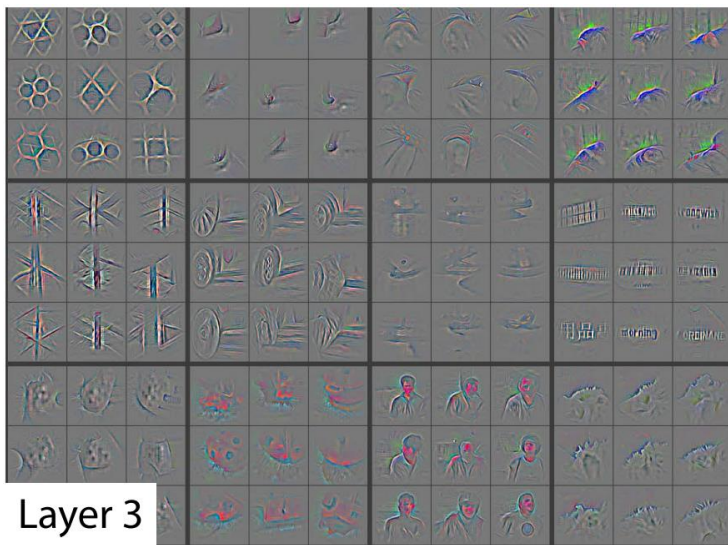

Intuition



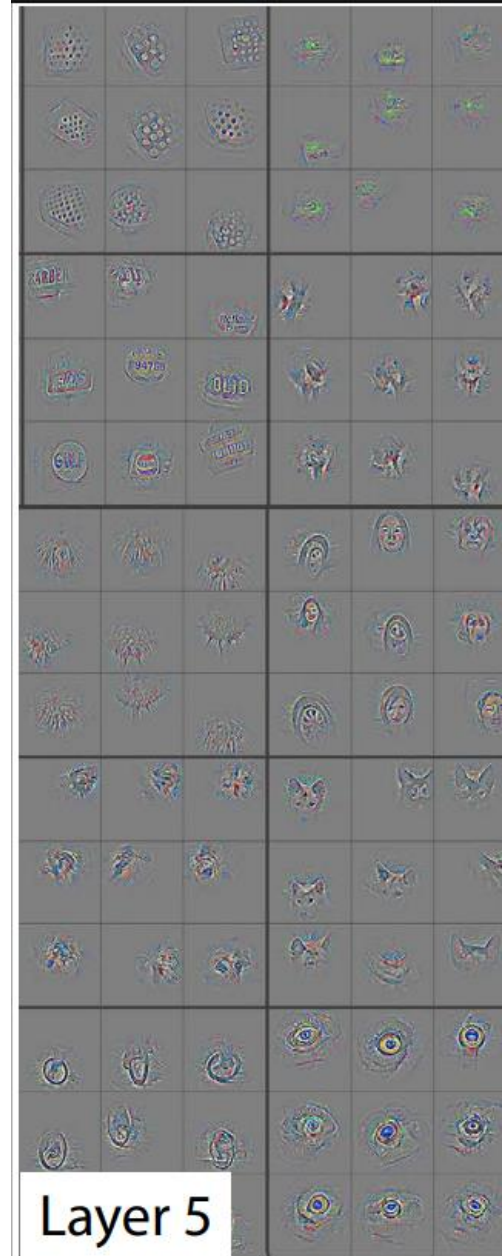

Layer 1



Layer 2



Layer 3




Layer 5




Max Pooling

- Reducing the spatial size of the feature map
- Example:
 - 2x2 max pooling with stride 2
 - For each 2x2 pixels in each channel, retain only the largest number
 - This type of pooling is extremely common in CNNs.



Max Pooling

- Pooling reduces the spatial dimension of the features
 - Not the number of channels
- Number of features is reduced by $2 \cdot 2$
 - Makes computation easier
 - A pooling layer is often followed convolution that doubles the number of features.
 - More higher level features, but lower resolution of them: good for classification



Types of pooling


- Max pooling
 - Max of the values
- Average pooling
 - Mean of the values
- Stochastic pooling
 - A random of the values
- Max pooling is the most common for classification

Max pooling – most used

- Max pooling is the most common for networks doing classification
 - For classification it does not matter much where exactly a feature is present
 - Taking the largest value helps make the model invariant to small translations
- Almost all poolings are 2x2 with stride=2.
 - Larger generate worse results.

Strided convolutions

- 3x3 convolution
Stride = 1




Input size: 9x9

Output size: 7x7

Strided convolutions

- 3x3 convolution
Stride = 2



Input size: 9x9

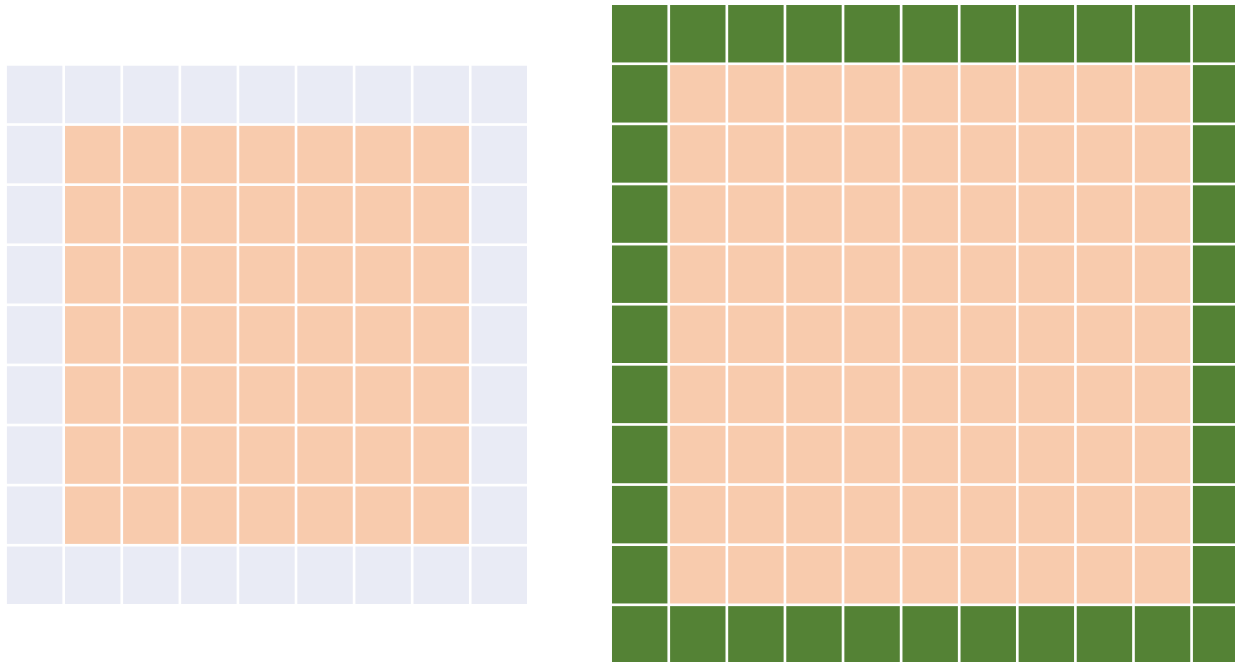
Output size: 4x4

Strided convolutions

- Can reduce the size of the image without using pooling
 - Computationally faster than pooling
 - At the risk of being less accurate
 - Preferred in some GANs due to consistent flow of gradients

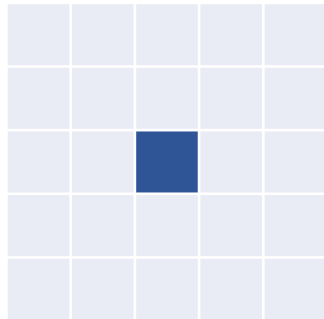
Padding

- The image size is reduced even for stride=1.
- Artificially increase the size of the image before convolution.
- Padding=1 will increase the size by one at the top, bottom, left and right.
- Usually zeros are used as the padding value in CNNs.




Padding

- Padding with zeros makes sense when the network contains feature maps
 - Positive values indicate the feature is present
 - Zero means this neuron was not activated (the activation was negative and ReLU made it zero)
- If the size needs to be kept constant, you need to pad with:
$$\frac{n-1}{2}$$
- Visualize the convolution kernel and see how much it has on each side
 - 5x5 \rightarrow padding=2




Question

- **DISCUSS WITH YOUR NEIGHBOUR 2 MINUTES**
- Input has 8 channels, and spatial dimensions 32x32
- We perform a 7x7 convolution that produces a new volume with 16 channels and still 32x32 spatially.
- What is the stride?
- What is the padding?
- How many weights (learnable parameters) does the convolution have?

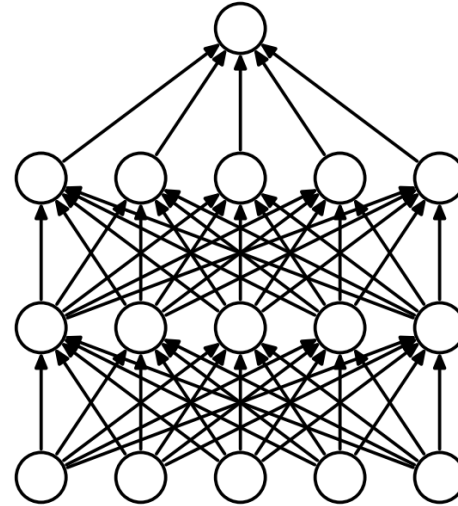


Question

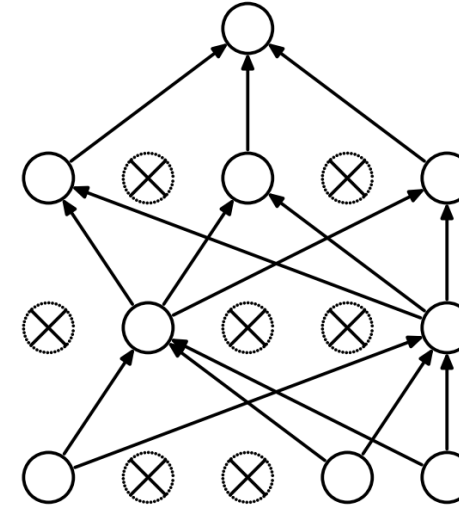
- What is the stride?
 - Stride=1
- What is the padding?
 - Padding=3
- How many weights (learnable parameters) does the convolution have?
 - $16 \times 8 \times 7 \times 7 + 16$ bias



Dropout



(a) Standard Neural Net



(b) After applying dropout.

- Background
 - Neural networks have many weights and can easily **overfit** to your data
- Concept: Model ensembles (averages of many) are always good
 - How can we do this in a single model?
- How it works
 - Each forward pass, we randomly omit each feature with a probability of 0.5
 - This means we are actually sampling from 2^n different architectures
 - Efficient way of performing model averaging with neural networks

Dropout


- Imagine if you each day only had either your right or left arm.
 - You would be forced to become good at using both arms, because you don't know which arm you will have tomorrow.
- Intuition:
 - Specialized neuron might be good enough to classify correctly
 - Other neurons become lazy
 - Randomly removing neurons forces all neurons to do their best

Dropout

- Technical details:
 - The dropped neurons during training means expectation of layer output is smaller than during training
 - This is problematic
 - It is common to scale the activations up by the dropout factor during training
 - For example if we drop $p=50\%$ of neurons, during training we multiply the activations by $1/p=2$
- This is often handled by the high-level framework you are using
 - As long as you tell the model whether it's training or testing right now.
- Dropout is a well known regularization technique, but BatchNormalization is an often used alternative.


Data augmentation

- Neural networks are very data-hungry
- How to satisfy their hunger?



Data augmentation



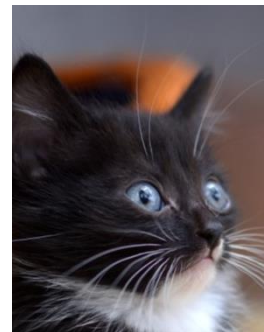
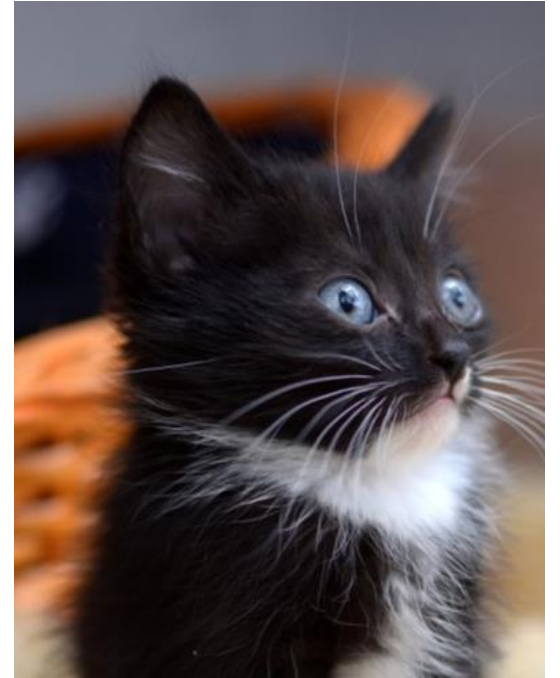
- Neural networks are very data-hungry
- How to satisfy their hunger?
- “just have enough data”
 - ImageNet has 14 million images
 - not big enough
 - Google has JFT-300M with 300 million images
 - not big enough
 - also not public
- We need to fully utilize the data we do have.



From: <https://ai.googleblog.com/2017/07/revisiting-unreasonable-effectiveness.html>

Data augmentation

- What is this?
- Any ideas on how we can create more images?



Data augmentation

- Create more data by exploiting things that the label is invariant towards
- Always keep in mind what the images are of
 - Flipping an image of text means that the text is no longer readable
- Types of augmentation:
 - Flips
 - Translation
 - Scale (non-uniform?)
 - Rotation
 - Elastic deformation
 - Lens distortion?
 - Noise
 - Blur
 - Small color changes

What you have learned

- Convolutions recap
- CNN
 - Convolutions
 - Max pooling
 - Stride/padding
 - Backprop
- Combatting overfitting:
 - Dropout
 - Data augmentation