

# Assignment-3 of Deep Learning in Computer Vision

## Generative Adversarial Networks

Xiao Hu<sup>1</sup> and Boya Wu<sup>2</sup>

1. xiahaa@space.dtu.dk, 2. s170061@student.dtu.dk

### 1 Generation of MNIST digits with a GAN

We have investigated a Generative Adversarial Networks (GAN) problem by generating digits in the MNIST-like fashion. The following tasks have been done:

- Implement a GAN as fully connected neural network and generate MNIST-like digits using Vanilla loss.
- Generate MNIST-like digits using least square loss.
- Implement DCGAN and generate MNIST-like digits.
- Implement cGAN and generate MNIST-like digits.
- Implement a cGAN that transform from SVHN to MNIST-like digits.


#### 1.1 GAN Design

For GAN, shown in Figure 1 :

- Generator: 4 fully connected layers with leaky ReLU (slope= 0.1) and batch normalization. Finally, the output is reshaped to a  $28 \times 28$  image and apply tanh to each pixel.
- Discriminator: 4 fully connected layers with hidden units being 1024, 512, 256, 1. Dropout is added after each fully connected layer.

For DCGAN, shown in Figure 2 :

- Generator: 4 deconvolution layers with leaky ReLU (slope= 0.1) and batch normalization followed by 1 convolution layer and use tanh to make each pixel ranging from  $-1$  to  $1$ . Feature channels are 1024, 512, 256, 128.
- Discriminator: 4 convolution layers with feature channels being 128, 256, 512, 1. The first three convolution layers are cascaded with a leaky ReLU (slope= 0.1) and dropout.



**Fig. 1.** GAN: discriminator (left) and Generator (right).

#### 1.2 Result

**GAN+Vanilla Loss** Result after 10 epochs training is shown in the left image of Figure 3.

**GAN+Least Square Loss** Result after 10 epochs training is shown in the middle image of Figure 3.




Fig. 2. DCGAN: discriminator (left) and Generator (right).




Fig. 3. Faked MNIST images generated by GAN+Vanilla loss (left), GAN+least square loss (middle), and DCGAN(right).

**DCGAN+Least Square Loss** Result after 10 epochs training is shown in the right image of Figure 3.

It can be seen DCGAN outperforms GAN. As a result, for the following parts, we keep using DCGAN and least square loss function.

**DCGAN+FashionMNIST** We applied the implemented DCGAN on FashionMNIST dataset. This time, the length of the random vector is increased to 300 since the dimension of the embedded manifold of FashionMNIST is higher than that of MNIST. Result after 10 epochs training is shown in the left image of Figure 4. It can be seen that the DCGAN is able to generate images in FashionMNIST style. However, compared with raw FashionMNIST images, it still misses some detailed texture. Training with more epochs or use higher dimension random vector may be helpful to get better result.


**cGAN** We use DCGAN to implement the conditional GAN. Result after 10 epochs training is shown in the right image of Figure 4. For each row, we draw the sample label and generated 10 samples.

**cGAN+SVHN** We use the implemented cGAN to transform from SVHN digits to MNIST digits. Result is shown in Figure 5. For each row, we randomly select 10 samples from SVHN dataset and then transform to MNIST style.

## 2 CycleGAN

The second task is to implement a CycleGAN to convert horses to zebras and vice versa. The following tasks have been done:

- Implement a CycleGAN.
- Train CycleGAN for image translation between horses and zebras.



**Fig. 4.** Faked FashionMNIST images generated by DCGAN+least square loss (left) and cGAN (right).

## 2.1 CycleGAN Network

The generator and discriminator for cycleGAN are shown in Figure 7. For building up CycleGAN, we actually need two generators:  $G_{A \rightarrow B}$ ,  $G_{B \rightarrow A}$ , and two discriminator:  $D_A$ ,  $D_B$ .

## 2.2 Loss Function

Three types of loss are needed:

- GAN loss: here we use least square GAN loss.
- Identity loss: e.g.  $|G_{A \rightarrow B}(a_{real}) - a_{real}|$ ,  $L_1$  norm is used here.
- Cycle loss: e.g.  $|G_{B \rightarrow A}(G_{A \rightarrow B}(a_{real})) - a_{real}|$ ,  $L_1$  norm is used here.

The final loss is the combination of the three loss:

$$loss_{total} = loss_{GAN} + \lambda_1 loss_{ind} + \lambda_2 loss_{cycle}$$

$\lambda_1 = 0.5$ ,  $\lambda_2 = 10$  are used here.

## 2.3 Training

Since training CycleGAN is really time consuming, we have not been able to train our network more than 10 epochs due to the time limitation. The result after 2 epochs is shown in We can see there are indeed some style tranfer effects. However, it is far from perfection. We think use more epochs may generate more promissing result.

## 2.4 Analysis

From the CycleGAN paper, right), there is an assumption that “We assume there is some underlying relationship between the domains for example, that they are two different renderings of the same underlying scene and seek to learn that relationship.”

## Appendix A Generator & Discriminator



Fig. 5. Style transfer from SVHN images to MNIST images.




Fig. 6. CycleGAN: discriminator (left) and Generator (right).

---

```

class ResNetBlock(nn.Module):
    def __init__(self, n_features):
        super(ResNetBlock, self).__init__()
        self.resblock = nn.Sequential(
            nn.Conv2d(n_features, n_features, kernel_size=3,
                      stride=1, padding=(1,1)),
            nn.InstanceNorm2d(n_features),
            nn.ReLU(),
            nn.Conv2d(n_features, n_features, kernel_size=3, stride=1, padding=(1,1))
        )

    def forward(self, x):
        out = self.resblock(x);
        out += x
        out = F.relu(out) ## could use leaky_relu?
        return out

```

---

```

def deconv(c_in, c_out, k_size, stride=2, padding=1, output_padding=0, bias = True,
           insn=True, relutype = 1, alpha=0.2):

```




Fig. 7. CycleGAN result after 2 epochs.

```

'''
return deconvolution layers
'''
layers = []
layers.append(nn.ConvTranspose2d(c_in, c_out, k_size, stride, padding, bias=bias,
                                output_padding=output_padding))

if relutype == 1:
    layers.append(nn.LeakyReLU(alpha))
else:
    layers.append(nn.ReLU())

if insn:
    layers.append(nn.InstanceNorm2d(c_out))
return nn.Sequential(*layers)

def conv(c_in, c_out, k_size, stride=2, padding=1, bias = True, insn=True, relutype = 1,
        alpha=0.2):
    '''
    return convolution layers
    '''
    layers = []
    layers.append(nn.Conv2d(c_in, c_out, k_size, stride, padding, bias=bias))
    if relutype == 1:
        layers.append(nn.LeakyReLU(alpha))
    else:
        layers.append(nn.ReLU())

    if insn:
        layers.append(nn.InstanceNorm2d(c_out))
    return nn.Sequential(*layers)

```

---

```

class G(nn.Module):
    '''
    generator
    '''
    def __init__(self, input_features = 3, output_features = 3):
        super(G, self).__init__()
        # encoding blocks
        features = 64
        self.conv1 = conv(input_features, features, 7, stride=1, padding=3, relutype=2)
        self.conv2 = conv(features, features*2, 3, stride=2, padding=1, relutype=2)
        self.conv3 = conv(features*2, features*4, 3, stride=2, padding=1, relutype=2) # 64 x
        64 x256

        # residual blocks
        num_res_blocks = 9

```

```

res_layers = []
for i in range(num_res_blocks):
    res_layers.append(ResNetBlock(features*4))
self.res_blocks = nn.Sequential(*res_layers)

# decoding blocks
self.deconv1 = deconv(features*4, features*2, 3, stride = 2, padding = 1,
    output_padding=1,relutype=2)
self.deconv2 = deconv(features*2, features, 3, stride = 2, padding = 1,
    output_padding=1,relutype=2)
self.conv4 = conv(features, output_features, 7, stride = 1, padding = 3,relutype=2)

def forward(self, x):
    out = self.conv1(x)    # (?, 64, 256, 256)
    out = self.conv2(out) # (?, 128, 128, 128)
    out = self.conv3(out)  # (?, 256, 64, 64)

    out = self.res_blocks(out)

    out = self.deconv1(out) # (?, 128, 128, 128)
    out = self.deconv2(out) # (?, 64, 256, 256) F.leaky_relu(
    out = torch.tanh(self.conv4(out))      # (?, 3, 256, 256)
    return out

```

---

```

class D(nn.Module):
    """Discriminator for mnist."""
    def __init__(self, input_features = 3, output_features = 1):
        super(D, self).__init__()
        conv_dim = 64
        self.conv1 = conv(input_features, conv_dim, 4, stride=2, padding=1, insn=False) ##
            64x128
        self.conv2 = conv(conv_dim, conv_dim*2, 4, stride=2, padding=1,insn=True) ## 128x64
        self.conv3 = conv(conv_dim*2, conv_dim*4, 4, stride=2, padding=1,insn=True) ##
            256x32
        self.conv4 = conv(conv_dim*4, conv_dim*8, 4, stride=1, padding=1,insn=True) ##
            512x31

        ## patch gan
        self.conv5 = conv(conv_dim*8, output_features, 4, stride=1, padding=1,insn=True) #
            1x30
        # self.conv6 = conv(output_features, output_features, 30, stride=1,
            padding=0) # 1x1

    def forward(self, x):
        out = self.conv1(x) # (?, 64, 128, 128)
        out = self.conv2(out) # (?, 128, 64, 64)
        out = self.conv3(out) # (?, 256, 32, 32)
        out = self.conv4(out) # (?, 512, 31, 31)
        out = self.conv5(out) # (?, 1, 30, 30)
        # out = self.conv6(out).squeeze() # (?,1,1,1)
        return out

```

---

## Appendix B Training

---

```

def lsgan_loss_d(dx, dz):
    b = 1
    return 0.5*torch.mean((dx - b)**2)+0.5*torch.mean((dz)**2)

```

```

def lsgan_loss_g(dz):
    c = 1
    return 0.5*torch.mean((dz-c)**2)

```

---

```

A_loader = horse_loader
B_loader = zebra_loader

beta1 = 0.5
beta2 = 0.9
lr = 0.0002

gA2B = G()
gB2A = G()
dA = D()
dB = D()

g_params = list(gA2B.parameters()) + list(gB2A.parameters())

g_optimizer = torch.optim.Adam(g_params, lr, [beta1, beta2])
da_optimizer = torch.optim.Adam(dA.parameters(), lr, [beta1, beta2])
db_optimizer = torch.optim.Adam(dB.parameters(), lr, [beta1, beta2])

gA2B.to(device)
gB2A.to(device)
dA.to(device)
dB.to(device)

A_iter = iter(A_loader)
B_iter = iter(B_loader)

```

---

```

iter_per_epoch = min(len(A_iter), len(B_iter))

# sampling image
fixed_A = A_iter.next()[0].to(device)
fixed_B = B_iter.next()[0].to(device)

visualization = ipywidgets.Output()
display.display(visualization)
with visualization:
    plt.figure(figsize=(20,20))
    subplots = [plt.subplot(1, 4, k+1) for k in range(4)]

criterion_cycle = torch.nn.L1Loss()
criterion_identity = torch.nn.L1Loss()

epochs = 30
for epoch in tqdm(range(epochs)):
    A_iter = iter(A_loader)
    B_iter = iter(B_loader)
    for step in tqdm(range(iter_per_epoch)):
        # load data
        imgA, label_A = A_iter.next()
        imgA, label_A = imgA.to(device), label_A.to(device)

        imgB, label_B = B_iter.next()
        imgB, label_B = imgB.to(device), label_B.to(device)

```

```

#===== train D =====#
# train with real images
d_optimizer.zero_grad()
outAA = dA(imgA)
with torch.no_grad():
    outBA = dA(gB2A(imgB))
da_gan_loss = lsgan_loss_d(outAA,outBA)

outBB = dB(imgB)
with torch.no_grad():
    outAB = dB(gA2B(imgA))
db_gan_loss = lsgan_loss_d(outBB,outAB)

da_gan_loss.backward()
db_gan_loss.backward()
d_optimizer.step()

#===== train G =====#
g_optimizer.zero_grad()

# Identity loss
loss_id_A = criterion_identity(gB2A(imgA), imgA)
loss_id_B = criterion_identity(gA2B(imgB), imgB)
loss_identity = (loss_id_A + loss_id_B) / 2

# add cycle loss
fake_B = gA2B(imgA)
g_gan_loss_A = lsgan_loss_g(dB(fake_B))
rec_A = gB2A(fake_B)
loss_cycle_A = criterion_cycle(rec_A, imgA)
#         g_cycle_loss = torch.mean((rec_A-imgA)**2)
fake_A = gB2A(imgB)
g_gan_loss_B = lsgan_loss_g(dA(fake_A))
rec_B = gA2B(fake_A)
loss_cycle_B = criterion_cycle(rec_B, imgB)

loss_cycle = (loss_cycle_A + loss_cycle_B) / 2
loss_GAN = (g_gan_loss_A+g_gan_loss_B)/2

# Total loss
g_loss = loss_GAN + 10 * loss_cycle + 0.5 * loss_identity

g_loss.backward()
g_optimizer.step()

# show the sampled images per 100
if (step+1) % 100 == 0:
    fake_A = gA2B(fixed_A)
    fake_B = gB2A(fixed_B)

    # show
    imga = fixed_A[0,:,:,:].cpu().numpy().transpose(1,2,0)
    imgfa = fake_A.cpu().detach()
    imgfa = imgfa[0,:,:,:].numpy()
    imgfa = imgfa.transpose(1,2,0)
    imga = imga * std + mean
    imgfa = imgfa * std + mean

    imgb = fixed_B[0,:,:,:].cpu().numpy().transpose(1,2,0)
    imgfb = fake_B.cpu().detach()
    imgfb = imgfb[0,:,:,:].numpy()
    imgfb = imgfb.transpose(1,2,0)

```



```

imgb = imgb * std + mean
imgfb = imgfb * std + mean

subplots[0].imshow(imga)
subplots[0].set_title('Raw A')
subplots[0].axis('off')

subplots[1].imshow(imgfa)
subplots[1].set_title('Fake A')
subplots[1].axis('off')

subplots[2].imshow(imgb)
subplots[2].set_title('Raw B')
subplots[2].axis('off')

subplots[3].imshow(imgfb)
subplots[3].set_title('Fake B')
subplots[3].axis('off')

display.display(plt.gcf())
display.clear_output(wait=True)

print('epoch: %d finish, save model'%(epoch))
model_path = './model'
#save the model parameters for each epoch
gAB_path = os.path.join(model_path, 'gab-%d.pkl'%(step+1))
gBA_path = os.path.join(model_path, 'gba-%d.pkl'%(step+1))
dA_path = os.path.join(model_path, 'da-%d.pkl'%(step+1))
dB_path = os.path.join(model_path, 'db-%d.pkl'%(step+1))
torch.save(self.gA2B.state_dict(), gAB_path)
torch.save(self.gB2A.state_dict(), gBA_path)
torch.save(self.dA.state_dict(), dA_path)
torch.save(self.dB.state_dict(), dB_path)

```

---